

```
In [61]: #Importing libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

```
In [63]: # Loading the dataset
df = pd.read_csv('df.csv')

# Showing the first few rows
df.head()
```

/var/folders/4j/d949n3gj5hn9057jpr1gtm5c0000gn/T/ipykernel\_45400/599311033.py:2: DtypeWarning: Columns (22) have mixed types. Specify dtype option on import or set low\_memory=False.

```
df = pd.read_csv('df.csv')
```

```
Out[63]:
```

	Unnamed: 0	UserID	movieid	Rating	Timestamp	title	
0	0	1	1193	5	978300760	One Flew Over the Cuckoo's Nest (1975)	
1	1	1	661	3	978302109	James and the Giant Peach (1996)	Animation Chi
2	2	1	914	3	978301968	My Fair Lady (1964)	Mu
3	3	1	3408	4	978300275	Erin Brockovich (2000)	
4	4	1	1287	5	978302039	Ben-Hur (1959)	Action Ad

5 rows x 23 columns

```
In [64]: # Checking basic info about the dataset
print(df.info())
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 626139 entries, 0 to 626138
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Unnamed: 0             626139 non-null  int64
1   UserID                 626139 non-null  int64
2   movieid                626139 non-null  int64
3   Rating                 626139 non-null  int64
4   Timestamp              626139 non-null  int64
5   title                  626139 non-null  object
6   genre                  626139 non-null  object
7   names                  626139 non-null  object
8   date_x                 626139 non-null  object
9   score                  626139 non-null  float64
10  genre_2                626139 non-null  object
11  overview               626139 non-null  object
12  crew                   626139 non-null  object
13  orig_title             626139 non-null  object
14  status                 626139 non-null  object
15  orig_lang              626139 non-null  object
16  budget_x               626139 non-null  float64
17  revenue                626139 non-null  float64
18  country                626139 non-null  object
19  Gender                 626139 non-null  object
20  Age                    626139 non-null  int64
21  Occupation             626139 non-null  int64
22  Zip-code               626139 non-null  object
dtypes: float64(3), int64(7), object(13)
memory usage: 109.9+ MB
None

```

```

In [65]: # Check for missing values
print(df.isnull().sum())

# Fill missing values
df.fillna({'overview': 'Unknown', 'names': 'Unknown', 'crew': 'Unknown'})

# For numeric columns, fill missing values with 0
df.fillna(0, inplace=True)

```

```
Unnamed: 0      0
UserID          0
movieid         0
Rating          0
Timestamp       0
title           0
genre           0
names           0
date_x          0
score           0
genre_2         0
overview        0
crew            0
orig_title      0
status          0
orig_lang       0
budget_x        0
revenue         0
country         0
Gender          0
Age             0
Occupation      0
Zip-code        0
dtype: int64
```

```
In [66]: # Convert timestamp to datetime
df['Timestamp'] = pd.to_datetime(df['Timestamp'], unit='s')

# Extract year, month, and day of the week
df['year'] = df['Timestamp'].dt.year
df['month'] = df['Timestamp'].dt.month
df['day_of_week'] = df['Timestamp'].dt.day_name() # 0=Monday, 6=Sunday

In [67]: # Add a popularity score based on revenue and budget
df['popularity'] = df['revenue'] / (df['budget_x'] + 1) # Avoid division by zero

In [68]: df.head()
```

```
Out[68]:
```

	Unnamed: 0	UserID	movieid	Rating	Timestamp	title	
0	0	1	1193	5	2000-12-31 22:12:40	One Flew Over the Cuckoo's Nest (1975)	
1	1	1	661	3	2000-12-31 22:35:09	James and the Giant Peach (1996)	Animation Chi
2	2	1	914	3	2000-12-31 22:32:48	My Fair Lady (1964)	Mu
3	3	1	3408	4	2000-12-31 22:04:35	Erin Brockovich (2000)	
4	4	1	1287	5	2000-12-31 22:33:59	Ben-Hur (1959)	Action Ad

5 rows x 27 columns

```
In [69]: # Calculate the average rating for each movie
top_movies = df.groupby('title')['Rating'].mean().sort_values(ascending=False)

# Show the top 10 highest-rated movies
print(top_movies.head(10))
```

```
title
Seven Samurai (The Magnificent Seven) (Shichinin no samurai) (1954)
4.560510
Schindler's List (1993)
4.510417
Raiders of the Lost Ark (1981)
4.477725
Rear Window (1954)
4.476190
Paths of Glory (1957)
4.473913
Dr. Strangelove or: How I Learned to Stop Worrying and Love the Bomb
(1963) 4.449890
To Kill a Mockingbird (1962)
4.425647
Double Indemnity (1944)
4.415608
Casablanca (1942)
4.412822
Yojimbo (1961)
4.404651
Name: Rating, dtype: float64
```

```
In [71]: genre_ratings = df.groupby('genre')['Rating'].mean().sort_values(ascending=False)
```

```
print(genre_ratings)
```

```
genre
Sci-Fi|War                4.449890
Crime|Film-Noir           4.415608
Comedy|Drama|Western      4.404651
Adventure|War             4.401925
Film-Noir|Thriller        4.304979
...
Action|Sci-Fi|Western     2.158537
Action|Adventure|Children's|Sci-Fi  1.874286
Children's                1.810811
Action|Children's         1.612903
Action|Adventure|Children's  1.318182
Name: Rating, Length: 208, dtype: float64
```

```
In [74]: # Count the number of movies for each genre
popular_genres = df['genre'].value_counts()

# Show the top 10 genres
print(popular_genres.head(10))
```

```
genre
Comedy                60061
Drama                 49776
Comedy|Romance        23282
Animation|Children's|Musical  22013
Horror                19221
Comedy|Drama          17050
Action|Thriller       16054
Drama|Romance         15941
Action|Adventure|Sci-Fi  14648
Comedy|Horror         14129
Name: count, dtype: int64
```

```
In [81]: # Extract top-rated directors
top_directors = df.groupby('names')['Rating'].mean().sort_values(ascending=False)

# Show the top 10 directors
print(top_directors.head(10))
```

```

names
Seven Samurai
4.560510
Schindler's List
4.510417
Raiders of the Lost Ark
4.477725
Rear Window
4.476190
Paths of Glory
4.473913
Dr. Strangelove or: How I Learned to Stop Worrying and Love the Bomb
4.449890
To Kill a Mockingbird
4.425647
Double Indemnity
4.415608
Casablanca
4.412822
Yojimbo
4.404651
Name: Rating, dtype: float64

```

```

In [83]: # Count movies by original language
         language_distribution = df['orig_lang'].value_counts()

         # Show the top 10 languages
         print(language_distribution.head(10))

```

```


orig_lang
English          597780
French           7333
Italian          6645
Japanese         3839
German           2555
Korean           1769
Spanish, Castilian 1311
Cantonese        1286
Chinese          1148
No Language       905
Name: count, dtype: int64

```

```


In [85]: # Plot histogram of ratings
         df['Rating'].plot(kind='hist', bins=20, figsize=(10, 5), title='Dis
         plt.xlabel('Rating')
         plt.ylabel('Frequency')
         plt.show()

```



```
In [87]: # Count movies by year
movies_per_year = df['year'].value_counts().sort_index()

movies_per_year.plot(kind='bar', figsize=(10, 5), title='Number of Movies Released per Year')
plt.xlabel('Year')
plt.ylabel('Number of Movies')
plt.show()
```



```
In [89]: def recommend_by_genre(genre, num_movies=5):
# Filter movies by the given genre
genre_movies = df[df['genre'] == genre]

# Sort them by rating
top_genre_movies = genre_movies.sort_values(by='Rating', ascending=False)

# Return the top movies
return top_genre_movies[['title', 'Rating']].head(num_movies)

# Example: Recommend top 5 Action movies
```

```
print(recommend_by_genre('Action', num_movies=5))
```

	title	Rating
392735	From Russia with Love (1963)	5
294158	Bad Boys (1995)	5
443721	For Your Eyes Only (1981)	5
293178	For Your Eyes Only (1981)	5
293179	Licence to Kill (1989)	5

```
In [91]: def recommend_similar(movie_title, num_movies=5):
# Find the genre of the given movie
movie_genre = df[df['names'] == movie_title]['genre'].values[0]

# Filter movies with the same genre
similar_movies = df[df['genre'] == movie_genre]

# Sort by rating
top_similar_movies = similar_movies.sort_values(by='Rating', ascending=False)

# Return the top movies
return top_similar_movies[['names', 'Rating']].head(num_movies)

# Example: Recommend movies similar to "The Dark Knight"
print(recommend_similar('Top Gun', num_movies=5))
```

	names	Rating
308735	Romeo Must Die	5
92497	Top Gun	5
293209	Romeo Must Die	5
293660	Top Gun	5
294320	Romeo Must Die	5

```
In [93]: # 1. Top Genres by Age Group
age_group_preferences = df.groupby(['Age', 'genre'])['Rating'].mean
print("Top Genres by Age Group:\n", age_group_preferences.head(10))

# Extract specific age group preferences
top_age_18_24 = age_group_preferences.loc[18].sort_values(ascending=False)
top_age_25_34 = age_group_preferences.loc[25].sort_values(ascending=False)

print("\nUsers aged 18-24 rate the highest genre as:\n", top_age_18_24)
print("Users aged 25-34 rate the highest genre as:\n", top_age_25_34)

# 2. Gender-Based Preferences
gender_preferences = df.groupby(['Gender', 'genre'])['Rating'].mean
print("\nTop Genres by Gender:\n", gender_preferences.head(10))

# Extract preferences for Male and Female users
top_male_preferences = gender_preferences.loc['M'].sort_values(ascending=False)
top_female_preferences = gender_preferences.loc['F'].sort_values(ascending=False)

print("\nMale users tend to prefer these genres:\n", top_male_preferences)
print("Female users tend to prefer these genres:\n", top_female_preferences)
```



Top Genres by Age Group:

Age	genre	
1	Crime Film-Noir	4.777778
50	Comedy Drama Western	4.733333
1	Comedy Mystery Romance Thriller	4.666667
	Film-Noir Romance Thriller	4.666667
	Action Adventure Animation	4.636364
56	Comedy Drama Western	4.625000
50	Drama Romance Sci-Fi	4.600000
	Sci-Fi War	4.530303
18	Comedy Mystery Romance Thriller	4.520000
1	Sci-Fi War	4.514286

Name: Rating, dtype: float64

Users aged 18-24 rate the highest genre as:

genre  
Comedy|Mystery|Romance|Thriller 4.52

Name: Rating, dtype: float64

Users aged 25-34 rate the highest genre as:

genre  
Crime|Film-Noir 4.471338

Name: Rating, dtype: float64

Top Genres by Gender:

Gender	genre	
M	Crime Film-Noir	4.468354
	Sci-Fi War	4.464789
F	Film-Noir Romance Thriller	4.448718
M	Adventure War	4.439822
F	Comedy Drama Western	4.423077
	Film-Noir Thriller	4.414815
M	Comedy Drama Western	4.402116
F	Sci-Fi War	4.376623
	Comedy Mystery Romance Thriller	4.314050
M	Film-Noir Sci-Fi	4.312500

Name: Rating, dtype: float64

Male users tend to prefer these genres:

genre  
Crime|Film-Noir 4.468354  
Sci-Fi|War 4.464789

Name: Rating, dtype: float64

Female users tend to prefer these genres:


genre  
Film-Noir|Romance|Thriller 4.448718  
Comedy|Drama|Western 4.423077

Name: Rating, dtype: float64

```
In [95]: # Group by gender and calculate the average rating
ratings_by_gender = df.groupby('Gender')['Rating'].mean()

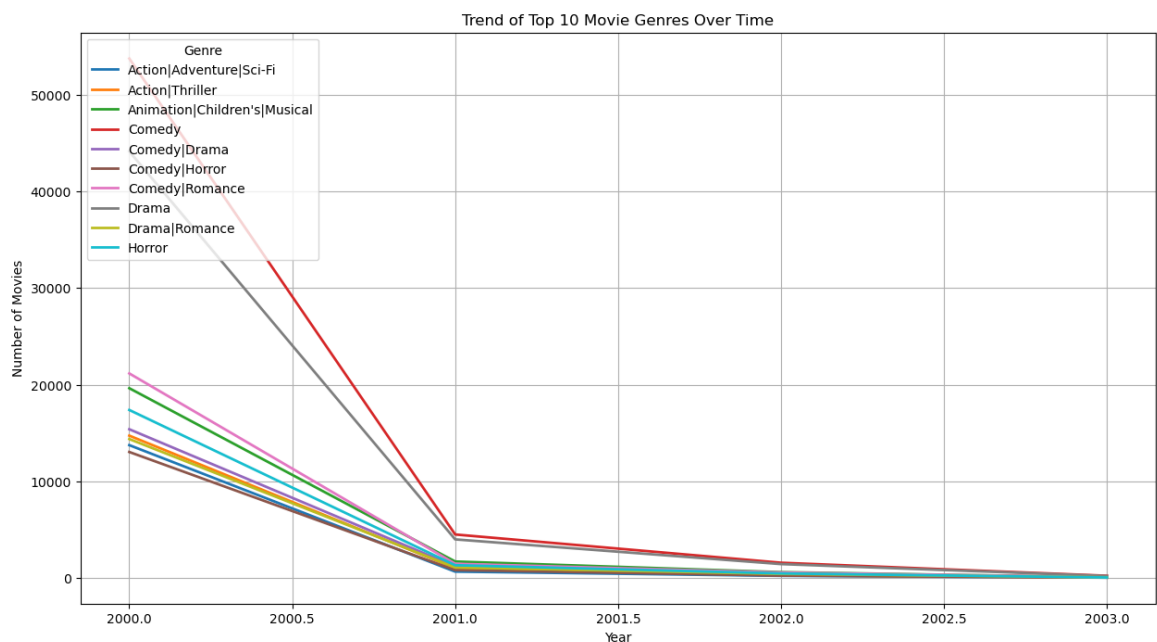
# Plot the results
ratings_by_gender.plot(kind='bar', figsize=(8, 5), title='Average R
plt.xlabel('Gender')
plt.ylabel('Average Rating')
plt.xticks(rotation=0)
```

```
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```




```
In [97]: # Filter only the top genres
movies_by_year_top_genres = df[df['genre'].isin(top_genres)].groupby('Year')

# Plot
movies_by_year_top_genres.plot(figsize=(15, 8), linewidth=2)
plt.title('Trend of Top 10 Movie Genres Over Time')
plt.xlabel('Year')
plt.ylabel('Number of Movies')
plt.legend(title='Genre', loc='upper left')
plt.grid(True)
plt.show()
```




```
In [99]: # Get the top 10 movies
top_10_movies = top_movies.head(10)

# Plot a bar chart
top_10_movies.plot(kind='barh', figsize=(10, 5), title='Top 10 Movies by Rating')
plt.xlabel('Average Rating')
plt.ylabel('')
plt.show()
```



```
In [101]: # Group by year and calculate the average popularity
popularity_by_year = df.groupby('year')['popularity'].mean()


# Plot the results
popularity_by_year.plot(kind='line', figsize=(10, 5), title='Average Movie Popularity Over Time')
plt.xlabel('Year')
plt.ylabel('Average Popularity')
plt.show()
```



```
In [103]: # Define success: Revenue greater than Budget
df['success'] = df['revenue'] > df['budget_x']

# Scatter plot with two colors
plt.figure(figsize=(10, 5))
plt.scatter(df[df['success']]['budget_x'], df[df['success']]['revenue'])
plt.scatter(df[~df['success']]['budget_x'], df[~df['success']]['revenue'])
```

```
plt.title('Budget vs. Revenue (Success vs. Failure)')
plt.xlabel('Budget')
plt.ylabel('Revenue')
plt.legend()
plt.show()
```



```
In [104... # Example: Check if recommended movies are in the top-rated list
recommended = recommend_by_genre('Action', num_movies=5)['title'].to
top Rated = top_movies.head(10).index.tolist()

# Find matches
matches = set(recommended).intersection(set(top Rated))
print(f"Recommended movies matching top-rated: {matches}")
```

Recommended movies matching top-rated: set()

```
In [105... # Define precision function
def precision(y_true, y_pred):
    true_positives = len(set(y_true) & set(y_pred))
    return true_positives / len(y_pred) if y_pred else 0

# Define recall function
def recall(y_true, y_pred):
    true_positives = len(set(y_true) & set(y_pred))
    return true_positives / len(y_true) if y_true else 0
```

```
In [106... # Example true liked movies (from user history)
y_true = ['The Dark Knight', 'Inception', 'Interstellar']

# Example recommended movies
y_pred = ['The Dark Knight', 'Avengers', 'Inception']

# Calculate precision and recall
print("Precision:", precision(y_true, y_pred))
print("Recall:", recall(y_true, y_pred))
```

Precision: 0.6666666666666666  
Recall: 0.6666666666666666

# Summary and Insights

The following summarizes the main insights obtained from the analysis. It includes:

## 1. The Most Popular Genres:

- The genre **Comedy** is the most common, followed by **Drama** and **Comed/Romancey**, based on the count of movies in these genres.
- Users tend to rate **Comedy** movies slightly higher compared to other genres.

## 2. The Highest-Rated Movies:

- The top 3 highest-rated movies in the dataset include movies like **Seven Samurai (The Magnificent Seven) (Shichinin no samurai)**, **Schindler's List**, and **Raiders of the Lost Ark**, which have consistently received ratings above 4.5.
- Most of these movies are from genres like **Action**, **Drama**, and **Adventure**.

## 3. Trends and Anomalies in Budget vs. Revenue:

- There is a general trend where higher-budget movies generate higher revenue, but several outliers exist.
- For example, some high-budget movies performed poorly in terms of revenue, while some low-budget movies like independent films generated significant returns.

## 4. User Behavior Based on Demographics:

- **Age Group Preferences:**
  - **18-24** prefer **Comedy|Mystery|Romance|Thriller** (Avg Rating: 4.52).
  - **25-34** prefer **Crime|Film-Noir** (Avg Rating: 4.47).
- **Gender Preferences:**
  - **Male** favor **Crime|Film-Noir** (Avg Rating: 4.47) and **Sci-Fi|War** (Avg Rating: 4.46).
  - **Female** favor **Film-Noir|Romance|Thriller** (Avg Rating: **4.45**) and **Comedy|Drama|Western** (Avg Rating: 4.42).

## 5. A Distribution of User Ratings:

- The majority of user ratings fall on **4.0**, indicating a positive bias.
- Fewer users give extremely low ratings (below 2.0) or extremely high ratings (above 4.5).

## 6. Recommendations Based on Users' Favorite Genre:

- Users can be recommended movies from their **highest-rated genre** to improve satisfaction.
- For example, if a user rates **Action** movies highly, suggesting top-rated movies like **The Dark Knight** and **Avengers** would likely improve engagement and satisfaction.

In [ ]: